

ALLaM: Large Language Models for Arabic and English

National Center for AI (NCAI), Saudi Data and AI Authority (SDAIA)

Riyadh, Saudi Arabia



1. Summary

The paper introduces ALLaM, a large language model designed to understand both Arabic and English. It was created to address the lack of advanced language models for Arabic, which has far fewer resources compared to English. ALLaM builds on existing models, like Llama-2, and improves them by carefully expanding the model's ability to learn Arabic while keeping its English skills intact. The authors used both natural Arabic text and machine-translated content, showing that this mix helps the model perform well in benchmarks like MMLU Arabic and ACVA. ALLaM also underwent fine-tuning to align it with human preferences, making it more useful and culturally aware.

2. Why This Paper?

I chose this paper because it highlights the importance of supporting the Arabic language in AI development. Many large language models focus primarily on English, leaving Arabic and other languages behind. ALLaM's approach is practical and sustainable, as it improves Arabic language understanding without starting from scratch. This work is inspiring for anyone interested in making technology more inclusive and accessible for Arabic-speaking communities.

3. The Limitations

As the authors mentioned, ALLaM was trained using internet data, which might include harmful language, unsafe content, and biases. This means the model could sometimes produce biased or toxic responses. Although ALLaM went through safety training to reduce these risks, it still needs more feedback from users to improve. Generative models like ALLaM can't be tested for every situation, so it's hard to predict how it will respond in all cases. This may lead to inaccurate, biased, or inappropriate outputs, even with neutral prompts. Developers should carefully test and fine-tune ALLaM for their specific needs. Finally, anything ALLaM generates doesn't represent the views of its creators or any related organization.

4. Benefiting From the Authors' Work

The work done in this paper can benefit many people. ALLaM sets a new standard for Arabic language models, helping improve applications in education, culture, and government services. And we saw a lot of great ideas in the recent ALLaM hackathon that SDAIA made in the past month. Developers can learn from ALLaM's techniques to adapt other AI models to low-resource languages, making AI more inclusive worldwide. Most importantly, this research shows how AI can support local languages and cultures while staying sustainable and efficient.